

# PATHspider II: The Tutorial

## Tutorial on Repeatability and Comparability in Measurement (RCM)

Iain R. Learmonth  
<iain@erg.abdn.ac.uk>

University of Aberdeen

ACM SIGCOMM 2018  
Monday, August 20, 2018

# Before we begin

- Please follow the instructions on your handout to get logged in to the test machine we will be using in this tutorial
- Each of you will have different instructions

# Active Measurement of Path Transparency

- Methodology:
  - 1 Throw packets at the Internet
  - 2 See what happens
- Ideal: two-ended A/B testing
- Scalable: one-ended A/B testing
- Multiple sources: **isolate** on-path from near-target impairment




Figure: One-Ended vs. Two-Ended Testing

# Introduction

## *PATHspider*

- The original implementation supported by mPlane/RITE
- Development continued in MAMI
- Three main functions:
  - Target list resolution
  - Active traffic generator
  - Passive observation
- Version 1.0 published at 2016 Applied Networking Research Workshop [1]




Figure: Original Architecture

# Introduction

## Active Traffic Generation

- Job queue created from input lists
- Connections made depending on the plugin
- Results will be merged with passive observation results




Figure: Original Architecture

# Introduction

## Built-In Flowmeter / Passive Observation

- PATHspider's built in flow meter is extensible.
- Using python-libtrace to dissect packets, any flow property imaginable can be reported back based on the raw packets:
  - ECN negotiation (IP/TCP headers)
  - Bleaching of bits, dropping of options
  - Checksum recalculations




Figure: Original Architecture

# Introduction

## Results from PATHspider 1.0


We presented some initial findings with the release of PATHspider:

### Explicit Congestion Notification (ECN)

State of ECN server-side deployment, as measured from a Digital Ocean vantage point in Amsterdam on 13th June 2016:

	IPv4	IPv6	all
<b>No ECN connectivity issues</b>	99.5%	99.9%	99.5%
<b>ECN successfully negotiated</b>	70.0%	82.8%	70.5%

ECN negotiation by Alexa rank bin:



### DiffServ Code Points (DSCP)

Initial study: 10,006 of 96,978 (10.31%) of Alexa Top 100k websites had unexpected, non-zero DSCP values. More measurement was needed to better characterize these anomalies.

### TCP Fast Open (TFO)

Initial study: 330 IPv4 and 32 IPv6 addresses of Alexa Top 1M are TFO-capable (of which 278 and 28 respectively are Google properties). DDoS prevention services, enterprise firewalls, and CPE tend to interfere with TFO. More measurement was necessary to analyze impairments.

# PATHspider 2.0

- Architecture changed to add a flow combiner
- Generalised to support more than just A/B testing
  - Any permutation of any number of tests
- Replaced PATHspider's HTTP code with cURL
- Added framework for packet forging based plugins using Scapy
- Completely rewritten (in Go) target list resolver
- Observer modules usable for standalone passive observation or analysis
- Source code: <https://github.com/mami-project/pathspider/tree/2.0.0/>




Figure: New Architecture



# Target Lists

## Types of Targets

- Popular - Cisco Umbrella, Alexa Topsites
- Curated Lists - CitizenLab
- Random - massscan

# Target Lists

*hellfire*

- Hellfire is a parallelised DNS resolver. It is written in Go and for the purpose of generating input lists to PATHspider, though may be useful for other applications
- Can use many sources for inputs:
  - Alexa Top 1 Million Global Sites
  - Cisco Umbrella 1 Million
  - Citizen Lab Test Lists
  - OpenDNS Public Domain Lists
  - Comma-Separated Values Files
  - Plain Text Domain Lists



# Target Lists

Using *hellfire*

```
1 $ hellfire
Usage:
3 hellfire --topsites [--file=<filename>] [--output=<individual|array|oneeach>] [--type=<
  host|ns|mx>] [--canid=<canid address>]
hellfire --cisco [--file=<filename>] [--output=<individual|array|oneeach>] [--type=<
  host|ns|mx>] [--canid=<canid address>]
5 hellfire --citizenlab [--country=<cc>|--file=<filename>] [--output=<individual|array|
  oneeach>] [--type=<host|ns|mx>] [--canid=<canid address>]
hellfire --opendns [--list=<name>|--file=<filename>] [--output=<individual|array|
  oneeach>] [--type=<host|ns|mx>] [--canid=<canid address>]
7 hellfire --csv --file=<filename> [--output=<individual|array|oneeach>] [--type=<host|ns
  |mx>] [--canid=<canid address>]
hellfire --txt --file=<filename> [--output=<individual|array|oneeach>] [--type=<host|ns
  |mx>] [--canid=<canid address>]
```

Listing 1: hellfire's Usage Help

```
$ hellfire --cisco
```

Listing 2: Start Resolving the Cisco Umbrella List

# Plugin Types

- Synchronised (traditional ecnspider)
  - ECN, DSCP
- Desynchronised (traditional ecnspider, no configurator)
  - TFO, H2, TLS NPN/ALPN
- Forge (new in PATHspider 2.0!)
  - Evil Bit, UDP Zero Checksum, UDP Options
- Single (new, and fast)
  - Various TCP Options

# Connection Helpers

- Instead of writing client code, use the code that already exists
- In the `pathspider.helpers` module:
  - DNS (`dnslib`)
  - HTTP/HTTPS (`pycURL`)
  - TCP (Python `socket`)
- For synchronised plugins, just use the helper
- For desynchronised plugins, the helpers are customisable, e.g. `cURL` helpers accept arbitrary `CURLOPTs`

# Synchronized Plugin

- SynchronizedSpider plugins use **built-in connection methods** along with **global system configuration** to change the behaviour of the connections
- Configuration functions are at the heart of a SynchronizedSpider plugin
- Configuration functions may make calls to sysctl or iptables to make changes to the way that traffic is generated.
- One function should be written for each of the configurations and PATHspider will ensure that the configurations are set before the corresponding traffic is generated. It is the responsibility of plugin authors to ensure that any configuration is reset by the next configuration function if that is required

# Synchronized Plugin

```
1 class ECN(SynchronizedSpider, PluggableSpider):
2     def config_no_ecn(self): # pylint: disable=no-self-use
3         """
4         Disables ECN negotiation via sysctl.
5         """
6
7         logger = logging.getLogger('ecn')
8         subprocess.check_call(
9             ['/sbin/sysctl', '-w', 'net.ipv4.tcp_ecn=2'],
10            stdout=subprocess.DEVNULL,
11            stderr=subprocess.DEVNULL)
12        logger.debug("Configurator disabled ECN")
13
14    def config_ecn(self): # pylint: disable=no-self-use
15        """
16        Enables ECN negotiation via sysctl.
17        """
18
19        logger = logging.getLogger('ecn')
20        subprocess.check_call(
21            ['/sbin/sysctl', '-w', 'net.ipv4.tcp_ecn=1'],
22            stdout=subprocess.DEVNULL,
23            stderr=subprocess.DEVNULL)
24        logger.debug("Configurator enabled ECN")
25
26    configurations = [config_no_ecn, config_ecn]
```

Listing 3: Configuration Functions for the ECN Plugin

# Explicit Congestion Notification

## Using the Built-In Plugin

```
$ sudo pspdr measure ecn < ~/2k-targets.ndjson > ~/ecn-results.ndjson
```

### Listing 4: Running the ECN plugin



# Desynchronized Plugin

- DesynchronizedSpider plugins modify the connection logic in order to change the behaviour of the connections. There is no global state synchronisation and so a DesynchronizedSpider can be more efficient than a SynchronizedSpider
- Connection functions are at the heart of a DesynchronizedSpider plugin
- These use a connection helper (or custom connection logic) to generate traffic towards with a target to get a reply from the target
- One function should be written for each connection to be made, usually with at least two functions to provide a baseline followed by an experimental connection

# Desynchronized Plugin

```
1 class H2(DesynchronizedSpider, PluggableSpider):
2     def conn_no_h2(self, job, config): # pylint: disable=unused-argument
3         curlopts = {}
4         curlinfos = [pycurl.INFO_HTTP_VERSION]
5         if self.args.connect == "http":
6             return connect_http(self.source, job, self.args.timeout, curlopts, curlinfos)
7         if self.args.connect == "https":
8             return connect_https(self.source, job, self.args.timeout, curlopts, curlinfos)
9         )
10        else:
11            raise RuntimeError("Unknown connection mode specified")
12
13    def conn_h2(self, job, config): # pylint: disable=unused-argument
14        curlopts = {pycurl.HTTP_VERSION: pycurl.CURL_HTTP_VERSION_2_0}
15        curlinfos = [pycurl.INFO_HTTP_VERSION]
16        if self.args.connect == "http":
17            return connect_http(self.source, job, self.args.timeout, curlopts, curlinfos)
18        if self.args.connect == "https":
19            return connect_https(self.source, job, self.args.timeout, curlopts, curlinfos)
20        )
21        else:
22            raise RuntimeError("Unknown connection mode specified")
23
24    connections = [conn_no_h2, conn_h2]
```

Listing 5: Connection Functions for the H2 Plugin

# Forge Plugin

- ForgeSpider plugins use Scapy to send forged packets to targets
- The heart of a ForgeSpider is the `forge()` function
- This function takes two arguments, the job containing the target information and the sequence number
- This function will be called the number of times set in the packets metadata variable and `seq` will be set to the number of times the function has been called for this job

# Single Plugin

- SingleSpider uses the built-in connection helpers to make a single connection to the target which is optionally observed by Observer chains
- This is the simplest model and only requires a `combine_flows()` function to generate conditions from the connection helper output and flow record output from the Observer

# Observer Modules

- While these used to be part of plugins in PATHspider 1.0, they are now independent and so can be reused across multiple plugins:
  - BasicChain, DNSChain, DSCPChain, ECNChain, EvilChain, ICMPChain, TCPChain, TFOChain
- These can also be used together, limiting each chain to just a single layer and letting the combiner produce conditions
- Chains can produce information to be consumed by other chains later in the list
- These can be used independently of a PATHspider measurement:

```
$ sudo pspdr observe tcp ecn
```

Listing 6: Running the PATHspider Observer independently

# Packet Forging

- PATHspider uses the Scapy library for Python for packet forging
- This is the most flexible method of creating new measurement plugins for PATHspider

# Make a Packet

- Scapy packets are constructed layer by layer
- While you can specify raw bytes, Scapy provides a number of useful classes for common protocols, which makes things a lot easier

# Make a Packet

- Scapy must be launched with `sudo` as we will need to use “raw” sockets to emit forged packets.

```
1 $ sudo scapy  
>>>
```

## Listing 7: Launching Scapy



# Make a Packet

## IPv4 Header - Create and Dissect

```
>>> IP()
2 <IP  |>
>>> i = IP()
4 >>> i.summary()
  '127.0.0.1 > 127.0.0.1 hopopt'
6 >>> i.display()
  ###[ IP ]###
8   version= 4
   ihl= None
10  tos= 0x0
   len= None
12  id= 1
   flags=
14  frag= 0
   ttl= 64
16  proto= hopopt
   chksum= None
18  src= 127.0.0.1
   dst= 127.0.0.1
20  \options\
```

Listing 8: Creating and Dissecting an IPv4 Header

# Make a Packet

## IPv4 Header - Customize

```
>>> i = IP(src="192.0.2.1",dst="198.51.100.1",ttl=10)
>>> i
<IP ttl=10 src=192.0.2.1 dst=198.51.100.1 |>
>>> i.summary()
'192.0.2.1 > 198.51.100.1 hopopt'
>>> i.display()
###[ IP ]###
version= 4
ihl= None
tos= 0x0
len= None
id= 1
flags=
frag= 0
ttl= 10
proto= hopopt
chksum= None
src= 192.0.2.1
dst= 198.51.100.1
\options\
```

Listing 9: Customizing an IPv4 Header

# Make a Packet

## TCP Header: Create and Dissect

```
>>> TCP()
2 <TCP |>
>>> t = TCP()
4 >>> t.summary()
   'TCP ftp_data > http S'
6 >>> t.display()
   ###[ TCP ]###
8   sport= ftp_data
   dport= http
10  seq= 0
   ack= 0
12  dataofs= None
   reserved= 0
14  flags= S
   window= 8192
16  chksum= None
   urgptr= 0
18  options= []
```

Listing 10: Creating and Dissecting a TCP Header

# Make a Packet

## TCP Header: Customizing

```
1 >>> t = TCP(dport=443)
2 >>> t
3 <TCP dport=https |>
4 >>> t.summary()
5 'TCP ftp_data > https S'
6 >>> t.display()
7 ###[ TCP ]###
   sport= ftp_data
   dport= https
   seq= 0
   ack= 0
   dataofs= None
   reserved= 0
   flags= S
   window= 8192
   chksum= None
   urgptr= 0
   options= []
```

Listing 11: Customizing a TCP Header

# Make a Packet

## Sticking the Pieces Together

- The / operator is used to join layers together.
- Scapy will automatically set fields, such as the IP Protocol field, when you do this.
- When dissecting, Scapy will automatically choose the dissector to use based on fields such as the IP Protocol field.

```
>>> p=i/t
>>> p.summary()
'IP / TCP 192.0.2.1:ftp_data > 198.51.100.1:https S'
>>> p.display()
[... output snipped ...]
```

Listing 12: Sticking the IP and TCP Headers Together

# Send a Packet

- The `sr1()` function sends a single packet, and returns a single packet if a reply is received.
- Start Wireshark capturing before executing the `sr1()` function.

```
1 >>> p=IP(dst="139.133.210.32")/TCP()
2 >>> a=sr1(p)
3 Begin emission:
4 . Finished sending 1 packets.
5 *
6 Received 2 packets, got 1 answers, remaining 0 packets
7 >>> a
<IP  version=4 ihl=5 tos=0x0 len=44 id=0 flags=DF frag=0 ttl=47 proto=tcp checksum=0xa98d
  src=139.133.210.32 dst=172.22.152.130 options=[] |<TCP  sport=http dport=ftp_data
  seq=3081101820 ack=1 dataofs=6 reserved=0 flags=SA window=29200 checksum=0xe9c7 urgptr
  =0 options=[('MSS', 1452)] |<Padding  load=':v' |>>>
9 >>> a.summary()
'IP / TCP 139.133.210.32: http > 172.22.152.130: ftp_data SA / Padding'
```

Listing 13: Create and Send an IP/TCP Packet

# Evil Bit

*The evil bit is a fictional **IPv4 packet header field** proposed in RFC 3514, a humorous April Fools' Day RFC from 2003 authored by Steve Bellovin. The RFC recommended that the last remaining unused bit, the "Reserved Bit," in the IPv4 packet header be used to indicate whether a packet had been sent with malicious intent, thus making computer security engineering an easy problem — simply ignore any messages with the evil bit set and trust the rest.*

– Wikipedia

# Evil Bit

## Setting the Evil Bit with Scapy

- The flags in the IP header are just an attribute you can modify:

```
>>> i = IP()  
>>> i.flags = 'evil'
```

Listing 14: Setting the Evil Bit on an IPv4 Header with Scapy



# PATHspider Plugins

## Directory Layout

- To get started you will need the required directory layout for PATHspider plugins, in this case for the EvilBit plugin:

```
pathspider-evilbit
+-- pathspider
    +-- __init__.py
    +-- plugins
        +-- __init__.py
        +-- evilbit.py
```

- Inside both `__init__.py` files, you will need to add the following (and only the following):

```
2 from pkgutil import extend_path
   __path__ = extend_path(__path__, __name__)
```

- Your plugin will be written in `evilbit.py` and this plugin will be discovered automatically when you run PATHspider

# PATHspider Plugins

## ForgeSpider

```
class EvilBit(ForgeSpider, PluggableSpider):  
    name = "evilbit"  
    description = "Evil bit connectivity testing"  
    version = '0.0.0'  
    chains = [BasicChain, TCPChain, EvilChain]  
    connect_supported = ["tcpsyn"]  
    packets = 2  
  
    def forge(self, job, seq):  
        ...
```

Listing 15: Outline for Evil Bit plugin using ForgeSpider

# PATHspider Plugins

## Forging the Packets

```
1 def forge(self, job, seq):  
    sport = 0  
3 while sport < 1024:  
    sport = int(RandShort())  
5 l4 = (TCP(sport=sport, dport=job['dp']))  
    if ':' in job['dip']:  
7 ip = IPv6(src=self.source[1], dst=job['dip'])  
    else:  
9 ip = IP(src=self.source[0], dst=job['dip'])  
    if seq == 1:  
11 ip.flags = 'evil'  
    return ip/l4
```

Listing 16: Creating Packets With and Without the Evil Bit

Up next:

## Path Transparency Observatory (PTO)

Don't delete your PATHspider results, you'll need them in the next session.

- Contact: [iain@erg.abdn.ac.uk](mailto:iain@erg.abdn.ac.uk)
- I review GitHub issues and pull requests regularly
- There is a `#pathspider` channel on `irc.oftc.net` which is where development discussion happens

# References I

- [1] Iain R. Learmonth, Brian Trammell, Mirja Kühlewind, and Gorrry Fairhurst.  
PATHspider: A tool for active measurement of path transparency.  
In *Proceedings of the 2016 Applied Networking Research Workshop*,  
pages 62–64, July 2016.